

# Import Modules

Import libraries required for ETL and Feature Selection

```
In [1]: # Set the seed value for the notebook so the results are reproducible
        from numpy.random import seed
        seed(42)

        import numpy as np
        import pandas as pd

        import warnings
        warnings.simplefilter('ignore', FutureWarning)

        %matplotlib inline
        import matplotlib.pyplot as plt

        from sklearn.model_selection import train_test_split
        from tensorflow.keras.utils import to_categorical
        from sklearn.preprocessing import LabelEncoder, StandardScaler, MinMaxScaler,
        RobustScaler, minmax_scale, MaxAbsScaler, Normalizer, QuantileTransformer, Pow
        erTransformer
```

## Method Definitions

### Plot Histogram Method

```
In [2]: # Code referenced from RealPython example: https://realpython.com/python-histograms/
        def plotHistogram(xValues, lbl):
            n, bins, patches = plt.hist(x=xValues, bins='auto', color='#0504aa', alpha
            =0.7, rwidth=0.85)
            plt.grid(axis='y', alpha=0.75)
            plt.xlabel(lbl)
            plt.ylabel('Frequency')
            plt.title(lbl + ' Histogram')
            maxfreq = n.max()

            # Set a clean upper y-axis limit.
            plt.ylim(ymax=np.ceil(maxfreq / 10) * 10 if maxfreq % 10 else maxfreq + 10
            )
```

```

In [3]: # Code to prep X and y
def prepData(data, scaler):

    # drop the dependent column from the independent columns
    X = data.drop(dependent, axis=1)

    # set the dependent column to y
    y = data[dependent]

    # Make a copy of the independent field matrix
    copyData = X.copy()

    # call pandas get dummies to convert to 1's and 0's in separate columns
    X = pd.get_dummies(copyData)

    # Split the data into test and train groups
    X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42,
stratify=y)

    # Save data out for use in other Jupyter Notebooks
    X_train_csv = X_train.to_csv ('X_train.csv', index = None, header=True)
    X_test_csv = X_test.to_csv ('X_test.csv', index = None, header=True)
    y_train_csv = y_train.to_csv ('y_train.csv', index = None, header=True)
    y_test_csv = y_test.to_csv ('y_test.csv', index = None, header=True)

    # Apply a scalar, first train then scale test and train data
    # TO DO: Insert function to select scalar
    X_scaler = scaler.fit(X_train)
    X_train_scaled = X_scaler.transform(X_train)
    X_test_scaled = X_scaler.transform(X_test)

    # Label encode the dependent variable, then Hot encode it
    label_encoder = LabelEncoder()
    label_encoder.fit(y_train)
    y_train_categorical = to_categorical(label_encoder.transform(y_train))
    y_test_categorical = to_categorical(label_encoder.transform(y_test))

    return X, y, X_train_scaled, X_test_scaled, y_train_categorical, y_test_categorical

```

## Arythmia - Cleveland Data Set

Create standard csv with header

```
In [4]: ##### Dependent variable (column) in our dataset #####
# Ayrthmia data set: We will use the binary sick/healthy column buff as our o
utput and remove the 'healthy' column
dependent = 'buff'
remove = 'healthy'

# Set the original file and new file names
originalFile = 'cleve.csv'
csvFile = 'cleveland.csv'
clean_csv = 'cleveland_clean.csv'

scalers = {"Robust": RobustScaler(),
           "Standard": StandardScaler(),
           "MinMax": MinMaxScaler(),
           "MaxAbs": MaxAbsScaler(),
           "Normalizer": Normalizer(),
           "Quantile": QuantileTransformer(),
           "Power": PowerTransformer()}
```

```
In [5]: # Open the Cleveland original file, add a header and convert to csv
with open(originalFile) as fin, open(csvFile, 'w') as fout:
    fout.write("age,sex,chestPain,trestBps,cholesterol,bloodSugar,ecg,maxHeart
Rate,exerciseInducedAngina,oldPeak,slope,vesselsColored,thal,buff,healthy\n")
    for line in fin:
        fout.write( ",".join(line.split()) + "\n" )
```

**Describe the data and look for unique values, remove bad data and any columns not needed**

```
In [6]: # Describe the continuous data and look for any anomalies
data = pd.read_csv(csvFile)
data.describe()
```

Out[6]:

	age	trestBps	cholesterol	maxHeartRate	oldPeak
count	303.000000	303.000000	303.000000	303.000000	303.000000
mean	54.366337	131.623762	246.264026	149.646865	1.039604
std	9.082101	17.538143	51.830751	22.905161	1.161075
min	29.000000	94.000000	126.000000	71.000000	0.000000
25%	47.500000	120.000000	211.000000	133.500000	0.000000
50%	55.000000	130.000000	240.000000	153.000000	0.800000
75%	61.000000	140.000000	274.500000	166.000000	1.600000
max	77.000000	200.000000	564.000000	202.000000	6.200000

```
In [7]: # determine column names to be able to look at all unique values for classification columns
data.columns
```

```
Out[7]: Index(['age', 'sex', 'chestPain', 'trestBps', 'cholesterol', 'bloodSugar',
              'ecg', 'maxHeartRate', 'exerciseInducedAngina', 'oldPeak', 'slope',
              'vesselsColored', 'thal', 'buff', 'healthy'],
              dtype='object')
```

```
In [8]: data['sex'].value_counts()
```

```
Out[8]: male    207
        fem     96
        Name: sex, dtype: int64
```

```
In [9]: data['chestPain'].value_counts()
```

```
Out[9]: asympt    143
        notang     87
        abnang     50
        angina     23
        Name: chestPain, dtype: int64
```

```
In [10]: data['bloodSugar'].value_counts()
```

```
Out[10]: fal     258
         true     45
         Name: bloodSugar, dtype: int64
```

```
In [11]: data['ecg'].value_counts()
```

```
Out[11]: norm    152
         hyp     147
         abn       4
         Name: ecg, dtype: int64
```

```
In [12]: data['exerciseInducedAngina'].value_counts()
```

```
Out[12]: fal     204
         true      99
         Name: exerciseInducedAngina, dtype: int64
```

```
In [13]: data['slope'].value_counts()
```

```
Out[13]: up       142
         flat     140
         down      21
         Name: slope, dtype: int64
```

```
In [14]: data['vesselsColored'].value_counts()
```

```
Out[14]: 0.0    175
          1.0     65
          2.0     38
          3.0     20
          ?       5
          Name: vesselsColored, dtype: int64
```

```
In [15]: # drop the data with rows where ? appear
          data.drop(data[data.vesselsColored == '?'].index, inplace=True)
```

```
In [16]: data['thal'].value_counts()
```

```
Out[16]: norm    163
          rev     115
          fix     18
          ?        2
          Name: thal, dtype: int64
```

```
In [17]: # drop the data with rows where ? appear
          data.drop(data[data.thal == '?'].index, inplace=True)
```

```
In [18]: data['buff'].value_counts()
```

```
Out[18]: buff    160
          sick    136
          Name: buff, dtype: int64
```

```
In [19]: data['healthy'].value_counts()
```

```
Out[19]: H      160
          S1      53
          S2      35
          S3      35
          S4      13
          Name: healthy, dtype: int64
```

```
In [20]: # We determined we were discarding this column and using the binary column 'buff' instead to indicate healthy vs sick
          # so delete this column
          data = data.drop(remove, axis=1)
```

```
In [21]: #display the data without ? rows
data
```

Out[21]:

	age	sex	chestPain	trestBps	cholesterol	bloodSugar	ecg	maxHeartRate	exercisIndu
0	63.0	male	angina	145.0	233.0	true	hyp	150.0	
1	67.0	male	asympt	160.0	286.0	false	hyp	108.0	
2	67.0	male	asympt	120.0	229.0	false	hyp	129.0	
3	37.0	male	notang	130.0	250.0	false	norm	187.0	
4	41.0	fem	abnang	130.0	204.0	false	hyp	172.0	
...	...	...	...	...	...	...	...	...	...
298	48.0	male	notang	124.0	255.0	true	norm	175.0	
299	57.0	male	asympt	132.0	207.0	false	norm	168.0	
300	49.0	male	notang	118.0	149.0	false	hyp	126.0	
301	74.0	fem	abnang	120.0	269.0	false	hyp	121.0	
302	54.0	fem	notang	160.0	201.0	false	norm	163.0	

296 rows × 14 columns


## Save our cleaned data set

```
In [22]: # overwrite the csv so we have a clean copy
export_csv = data.to_csv(clean_csv, index = None, header=True)
```


## Now plot each continuous column in a histogram to examine the distribution

These can be referenced when selecting machine learning options

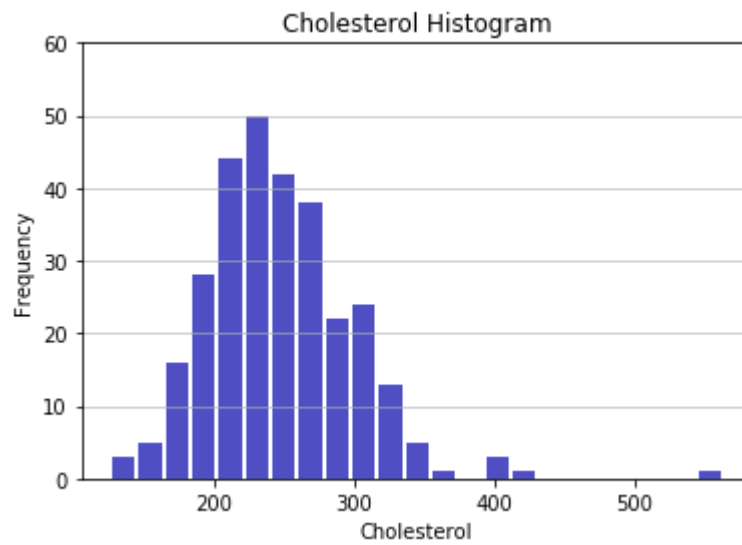
```
In [23]: plotHistogram(data['age'], 'Age')
```



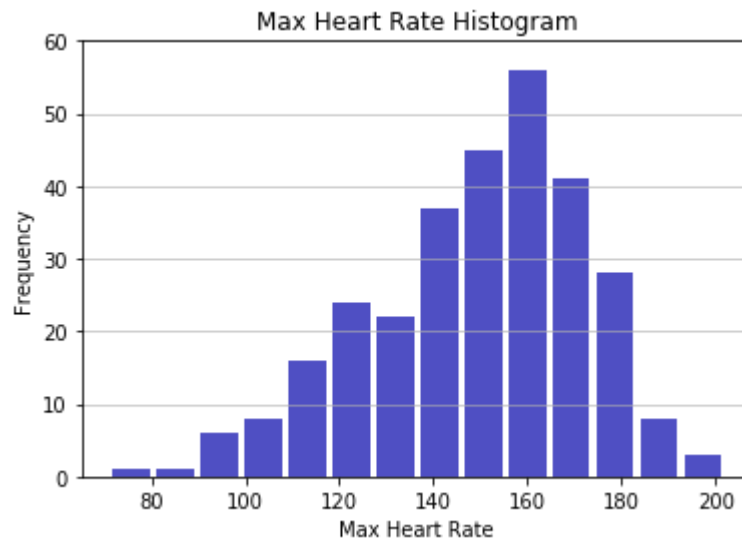
```
In [24]: plotHistogram(data['trestBps'], 'Trest Bps')
```



```
In [25]: plotHistogram(data['cholesterol'], 'Cholesterol')
```




```
In [26]: plotHistogram(data['maxHeartRate'], 'Max Heart Rate')
```





```
In [27]: plotHistogram(data['oldPeak'], 'Old Peak')
```



## Let's Try Some Feature Selection Techniques

This is an attempt to limit the number of features used in the models to simplify and remove noise

**NOTE: These three methods and code were obtained from the following website:**

<https://towardsdatascience.com/feature-selection-techniques-in-machine-learning-with-python-f24e7da3f36e>  
(<https://towardsdatascience.com/feature-selection-techniques-in-machine-learning-with-python-f24e7da3f36e>)

**First we need to create our X and y**

```
In [28]: # X, y, X_train_scaled, X_test_scaled, y_train_categorical, y_test_categorical
= prepData(data, scalers['MinMax'])
X, y, X_train_scaled, X_test_scaled, y_train_categorical, y_test_categorical =
prepData(data, scalers['MinMax'])
X
```

Out[28]:

	age	trestBps	cholesterol	maxHeartRate	oldPeak	sex_fem	sex_male	chestPain_abnang
0	63.0	145.0	233.0	150.0	2.3	0	1	0
1	67.0	160.0	286.0	108.0	1.5	0	1	0
2	67.0	120.0	229.0	129.0	2.6	0	1	0
3	37.0	130.0	250.0	187.0	3.5	0	1	0
4	41.0	130.0	204.0	172.0	1.4	1	0	1
...	...	...	...	...	...	...	...	...
298	48.0	124.0	255.0	175.0	0.0	0	1	0
299	57.0	132.0	207.0	168.0	0.0	0	1	0
300	49.0	118.0	149.0	126.0	0.8	0	1	0
301	74.0	120.0	269.0	121.0	0.2	1	0	1
302	54.0	160.0	201.0	163.0	0.0	1	0	0

296 rows × 28 columns

## Univariate Selection using ChiSquare

```
In [29]: # REference: https://towardsdatascience.com/feature-selection-techniques-in-machine-learning-with-python-f24e7da3f36e
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2

#apply SelectKBest class to extract best features
bestfeatures = SelectKBest(score_func=chi2, k=10)
fit = bestfeatures.fit(X_train_scaled,y_train_categorical)
dfscores = pd.DataFrame(fit.scores_)
dfcolumns = pd.DataFrame(X.columns)

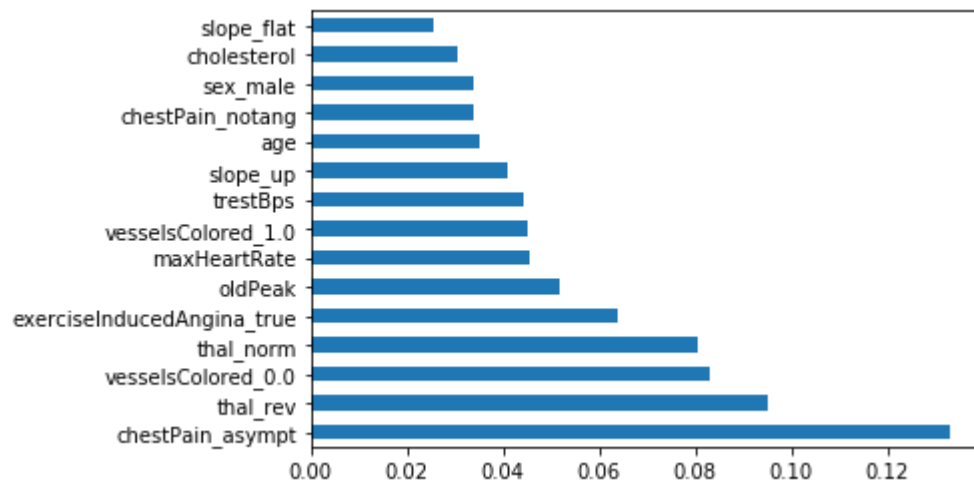
#concat two dataframes for better visualization
featureScores = pd.concat([dfcolumns,dfscores],axis=1)
featureScores.columns = ['Specs', 'Score'] #naming the dataframe columns
print(featureScores.nlargest(27,'Score')) #print best features reverse order
```

	Specs	Score
27	thal_rev	34.200294
17	exerciseInducedAngina_true	33.418728
9	chestPain_asympt	32.604218
26	thal_norm	28.482451
21	vesselsColored_0.0	21.540903
20	slope_up	18.872001
19	slope_flat	15.823077
16	exerciseInducedAngina_fal	14.756322
7	chestPain_abnang	14.126050
10	chestPain_notang	13.687783
23	vesselsColored_2.0	12.695136
5	sex_fem	12.124316
22	vesselsColored_1.0	11.352036
4	oldPeak	10.745139
24	vesselsColored_3.0	7.004118
6	sex_male	5.700838
13	ecg_abn	3.529412
8	chestPain_angina	2.826471
15	ecg_norm	2.605742
3	maxHeartRate	1.955725
14	ecg_hyp	1.548633
0	age	0.975826
25	thal_fix	0.741422
18	slope_down	0.684007
1	trestBps	0.355538
2	cholesterol	0.136939
12	bloodSugar_true	0.022368

## Feature Importance

```
In [30]: from sklearn.ensemble import ExtraTreesClassifier
import matplotlib.pyplot as plt
model = ExtraTreesClassifier()
model.fit(X_train_scaled,y_train_categorical)
print(model.feature_importances_) #use inbuilt class feature_importances of tree based classifiers
#plot graph of feature importances for better visualization
feat_importances = pd.Series(model.feature_importances_, index=X.columns)
feat_importances.nlargest(15).plot(kind='barh')
plt.show()
```

```
[0.03502496 0.04442578 0.03073113 0.04561243 0.0516222 0.01905353
0.03392827 0.00337018 0.0093416 0.13287963 0.03397756 0.0074519
0.01069452 0.00024183 0.02261995 0.02012887 0.01899636 0.06371845
0.01080309 0.02556698 0.04095077 0.08324723 0.04526443 0.0121062
0.01591545 0.00653634 0.08058911 0.09520124]
```



## Correlation Matrix with Heatmap

Notice that only continuous variables are plotted and there are no variables with a high correlation in our dataset


```
In [31]: import seaborn as sns
```

```
#get correlations of each features in dataset
corrmat = data.corr()
top_corr_features = corrmat.index
plt.figure(figsize=(5,5))

#plot heat map
g=sns.heatmap(data[top_corr_features].corr(),annot=True,cmap="RdYlGn")
corrmat
```

Out[31]:

	age	trestBps	cholesterol	maxHeartRate	oldPeak
age	1.000000	0.288701	0.200920	-0.395989	0.199373
trestBps	0.288701	1.000000	0.126705	-0.051817	0.196793
cholesterol	0.200920	0.126705	1.000000	-0.001947	0.042144
maxHeartRate	-0.395989	-0.051817	-0.001947	1.000000	-0.346742
oldPeak	0.199373	0.196793	0.042144	-0.346742	1.000000



In [ ]: